



# DS-GA 1007

## Programming for Data Science

Lecture 10

pandas I - Operations on Tables



Package that combines  
array operations and  
queries on tabular data

# DS-GA 1007

## Programming for Data Science

### Lecture 10

### pandas I - Operations on Tables

# Announcements

- ▶ Homework 7 due **Sunday November 10** at 11:59pm
- ▶ Survey 3 due **Sunday November 10** at 11:59pm
- ▶ Project
  - ▶ Milestone due **Thursday November 28** at 11:59pm
  - ▶ Background
  - ▶ Plans
    - ▶ Some Components of the Software
    - ▶ Some Relevant Datasets and Approaches



# Announcements

- ▶ Homework 7 due **Sunday November 10** at 11:59pm
- ▶ Survey 3 due **Sunday November 10** at 11:59pm
- ▶ Project
  - ▶ Milestone due **Thursday November 28** at 11:59pm
  - ▶ Background
  - ▶ Plans
    - ▶ Some Components of the Software
    - ▶ Some Relevant Datasets and Approaches

How was lab on Monday?!



# Agenda

- ▶ Review

  - ▶ numpy

  - ▶ matplotlib

- ▶ Lesson

  - ▶ notebooks and scripts

  - ▶ pandas

- ▶ Readings

  - ▶ Python for Data Analysis by Wes McKinney

  - ▶ <http://pandas.pydata.org/pandas-docs/stable/index.html>

## Objectives

- ▶ How can we perform operations on arrays?

- ▶ How can we choose appropriate charts for data types?

- ▶ Is it possible to link notebook and scripts?

- ▶ How can we arrange rows and columns of a table?

# Agenda

- ▶ Review



numpy

- ▶ matplotlib

- ▶ Lesson

- ▶ notebooks and scripts

- ▶ pandas

- ▶ Readings

- ▶ Python for Data Analysis by Wes McKinney

- ▶ <http://pandas.pydata.org/pandas-docs/stable/index.html>

## Objectives

- ▶ Dimensions, Shape and Data Type
- ▶ Retrieving Data
  - ▶ Access, Iterate
  - ▶ Slice, Mask
- ▶ Broadcasting
- ▶ Views vs Copies

# Agenda

- ▶ Review



numpy

- ▶ matplotlib

- ▶ Lesson

- ▶ notebooks and scripts

- ▶ pandas

- ▶ Readings

- ▶ Python for Data Analysis by Wes McKinney

- ▶ <http://pandas.pydata.org/pandas-docs/stable/index.html>

## Objectives

- ▶ Dimensions, Shape and Data Type
- ▶ Retrieving Data
  - ▶ Access, Iterate
  - ▶ Slice, Mask
- ▶ Broadcasting
- ▶ Views vs Copies

```
import numpy  
import numpy as np  
from numpy import *
```

# Dimension, Shape and Data Type

- ▶ Dimension
  - ▶ Number of nested arrays
  - ▶ Each array dimension called *axis*
  - ▶ Axes start at 0
  - ▶ Number of Axes is *rank*
- ▶ *ndarray.ndim*

		axis 1		
		0	1	2
axis 0	0	(0,0) (-3,-3)	(0,1) (-3,-2)	(0,2) (-3,-1)
	1	(1,0) (-2,-3)	(1,1) (-2,-2)	(1,2) (-2,-1)
	2	(2,0) (-1,-3)	(2,1) (-1,-2)	(2,2) (-1,-1)

rank = 2



# Dimension, Shape and Data Type

## ► Shape

- Tuple of integers indicating size of each axis
- **Size** of array is total number of elements
- *ndarray.shape*
- *ndarray.size*

		axis 1		
		0	1	2
axis 0	0	(0,0) (-3,-3)	(0,1) (-3,-2)	(0,2) (-3,-1)
	1	(1,0) (-2,-3)	(1,1) (-2,-2)	(1,2) (-2,-1)
	2	(2,0) (-1,-3)	(2,1) (-1,-2)	(2,2) (-1,-1)

rank = 2

# Dimension, Shape and Data Type

## ► Data Type

- The package will try to detect the data type from the entries
- Many data types and precision supported
  - float32
  - int8
  - complex64

## ► *ndarray.dtype*

		axis 1		
		0	1	2
axis 0	0	(0,0) (-3,-3)	(0,1) (-3,-2)	(0,2) (-3,-1)
	1	(1,0) (-2,-3)	(1,1) (-2,-2)	(1,2) (-2,-1)
	2	(2,0) (-1,-3)	(2,1) (-1,-2)	(2,2) (-1,-1)

**rank = 2**

# Retrieving Data

## ► Access

- Arrays indexed by tuple of integers
- Use brackets to get entries
- Negative indices are supported
- $x[i,j]$  or  $x[(i,j)]$

		axis 1		
		0	1	2
axis 0	0	(0,0) (-3,-3)	(0,1) (-3,-2)	(0,2) (-3,-1)
	1	(1,0) (-2,-3)	(1,1) (-2,-2)	(1,2) (-2,-1)
	2	(2,0) (-1,-3)	(2,1) (-1,-2)	(2,2) (-1,-1)

rank = 2

# Retrieving Data

## ► Iterate

- In loops iteration performed over first axes
- Can convert from two dimensional to one dimensional in row major or column major order

```
a2d = np.array([[1,3,5,7,9,11],  
               [2,4,6,8,10,12],  
               [0,1,2,3,4,5]])
```

```
for r in a2d:  
    print(r)
```

```
[ 1  3  5  7  9 11]  
[ 2  4  6  8 10 12]  
[ 0  1  2  3  4  5]
```

```
for i in a2d.flat:  
    print(i)
```

```
1  
3  
5  
7  
9
```

# Retrieving Data

- ▶ Slicing
  - ▶ Specified by start, end and stride
  - ▶ Omitting an index means from the beginning, to the end or default stride
  - ▶ A slice is a **view** of the original array. Views are like references

a2d

```
array([[ 1,  3,  5,  7,  9, 11],  
       [ 2,  4,  6,  8, 10, 12],  
       [ 0,  1,  2,  3,  4,  5]])
```

```
print('fixing a column and traversing rows\n',a2d[:,2])
```

```
print('traversing an array block\n',a2d[1:,2:5])
```

```
print('traversing a subset of rows \n',a2d[[0,2]])
```

```
print('traversing a subset of columns \n',a2d[:,[0,2,5]])
```

```
print('traversing a subset of array elements \n',a2d[[0,1,2],[0,2,5]])
```

# Retrieving Data

- ▶ Slicing
  - ▶ Specified by start, end and stride
  - ▶ Omitting an index means from the beginning, to the end or default stride
  - ▶ A slice is a **view** of the original array. Views are like references
- ▶ Fancy Indexing
  - ▶ Passing multiple pairs of indices

a2d

```
array([[ 1,  3,  5,  7,  9, 11],  
       [ 2,  4,  6,  8, 10, 12],  
       [ 0,  1,  2,  3,  4,  5]])
```

```
print('fixing a column and traversing rows\n',a2d[:,2])
```

```
print('traversing an array block\n',a2d[1:,2:5])
```

```
print('traversing a subset of rows \n',a2d[[0,2]])
```

```
print('traversing a subset of columns \n',a2d[:,[0,2,5]])
```

```
print('traversing a subset of array elements \n',a2d[[0,1,2],[0,2,5]])
```

# Retrieving Data

- ▶ Slicing
  - ▶ Specified by start, end and stride
  - ▶ Omitting an index means from the beginning, to the end or default stride
  - ▶ A slice is a **view** of the original array. Views are like references
- ▶ Fancy Indexing
  - ▶ Passing multiple pairs of indices

a2d

```
array([[ 1,  3,  5,  7,  9, 11],  
       [ 2,  4,  6,  8, 10, 12],  
       [ 0,  1,  2,  3,  4,  5]])
```

```
print('fixing a column and traversing rows\n',a2d[:,2])
```

```
print('traversing an array block\n',a2d[1:,2:5])
```

```
print('traversing a subset of rows \n',a2d[[0,2]])
```

```
print('traversing a subset of columns \n',a2d[:,[0,2,5]])
```

```
print('traversing a subset of array elements \n',a2d[[0,1,2],[0,2,5]])
```

traversing an array block

```
[[ 6  8 10]
```

```
 [ 2  3  4]]
```

traversing a subset of rows

```
[[ 1  3  5  7  9 11]
```

```
 [ 0  1  2  3  4  5]]
```

traversing a subset of columns

```
[[ 1  5 11]
```

```
 [ 2  6 12]
```

```
 [ 0  2  5]]
```

traversing a subset of array elements

```
[1 6 5]
```

# Retrieving Data

## ► Masking

- Array of True and False useful to select elements in another array
- Can make assignments using True and False

```
x = np.arange(18).reshape(3,6)
print(x)
mask = (x > 7)
print(mask)
print(x[mask])
x[mask]=0
print(x)
```

```
[[ 0  1  2  3  4  5]
 [ 6  7  8  9 10 11]
 [12 13 14 15 16 17]]
[[False False False False False False]
 [False False  True  True  True  True]
 [ True  True  True  True  True  True]]
[ 8  9 10 11 12 13 14 15 16 17]
[[0 1 2 3 4 5]
 [6 7 0 0 0 0]
 [0 0 0 0 0 0]]
```



# Views vs Copies

- ▶ Slicing an array gives a view
  - ▶ Like a reference to part to an array
  - ▶ Changing elements of the view will change the original array
- ▶ Accessing an array through indexing gives a copy
- ▶ Use *copy.deepcopy* to avoid side-effects

```
x = np.arange(18).reshape(3,6)
print(x)
mask = (x > 7)
print(mask)
print(x[mask])
x[mask]=0
print(x)
```

```
[[ 0  1  2  3  4  5]
 [ 6  7  8  9 10 11]
 [12 13 14 15 16 17]]
[[False False False False False False]
 [False False  True  True  True  True]
 [ True  True  True  True  True  True]]
[ 8  9 10 11 12 13 14 15 16 17]
[[0 1 2 3 4 5]
 [6 7 0 0 0 0]
 [0 0 0 0 0 0]]
```

# Broadcasting

- ▶ Allows for element by element operations on arrays with different sizes
- ▶ Package transforms arrays so that they all have the same size

```
A = np.arange(25).reshape(5,5)
s = 5
B = s + A
B
```

```
array([[ 5,  6,  7,  8,  9],
       [10, 11, 12, 13, 14],
       [15, 16, 17, 18, 19],
       [20, 21, 22, 23, 24],
       [25, 26, 27, 28, 29]])
```

$$B = s + A = \begin{bmatrix} 5 & 5 & 5 & 5 & 5 \\ 5 & 5 & 5 & 5 & 5 \\ 5 & 5 & 5 & 5 & 5 \\ 5 & 5 & 5 & 5 & 5 \\ 5 & 5 & 5 & 5 & 5 \end{bmatrix} + \begin{bmatrix} 0 & 1 & 2 & 3 & 4 \\ 5 & 6 & 7 & 8 & 9 \\ 10 & 11 & 12 & 13 & 14 \\ 15 & 16 & 17 & 18 & 19 \\ 20 & 21 & 22 & 23 & 24 \end{bmatrix}$$

# Broadcasting

- ▶ The two arrays must be compatible for broadcasting
  - ▶ the axis lengths match, or
  - ▶ either of the lengths is 1
- ▶ Broadcasting is then performed over the missing and/or length 1 dimensions

```
v = np.arange(5)
print(v.shape)

w = np.arange(3).reshape(3,1)
print(w.shape)

Z = v + w

print(v, '\n')
print(w, '\n')
print(Z)
```

```
(5,)
(3, 1)
[0 1 2 3 4]
```

```
[[0]
 [1]
 [2]]
```

```
[[0 1 2 3 4]
 [1 2 3 4 5]
 [2 3 4 5 6]]
```

# Broadcasting

- ▶ The two arrays must be compatible for broadcasting
  - ▶ the axis lengths match, or
  - ▶ either of the lengths is 1
- ▶ Broadcasting is then performed over the missing and/or length 1 dimensions

$$Z = v + w$$

↓

$$[0 \quad 1 \quad 2 \quad 3 \quad 4] + \begin{bmatrix} 0 \\ 1 \\ 2 \end{bmatrix}$$

↓

$$\begin{bmatrix} 0 & 1 & 2 & 3 & 4 \\ 0 & 1 & 2 & 3 & 4 \\ 0 & 1 & 2 & 3 & 4 \end{bmatrix} + \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 1 & 1 & 1 \\ 2 & 2 & 2 & 2 & 2 \end{bmatrix}$$

# Vectorized Operations

- ▶ Repeated operations in loops are slow
- ▶ Package supports C extensions that allow for execution across arrays element by element
- ▶ Processing concurrently instead of sequentially saves time

```
import numpy as np
np.random.seed(0)

def compute_reciprocals(values):
    output = np.empty(len(values))
    for i in range(len(values)):
        output[i] = 1.0 / values[i]
    return output

values = np.random.randint(1, 10, size=5)
compute_reciprocals(values)
```

# Vectorized Operations

- ▶ Repeated operations in loops are slow
- ▶ Package supports C extensions that allow for execution across arrays element by element
- ▶ Processing concurrently instead of sequentially saves time

```
import numpy as np
np.random.seed(0)

def compute_reciprocals(values):
    output = np.empty(len(values))
    for i in range(len(values)):
        output[i] = 1.0 / values[i]
    return output

values = np.random.randint(1, 10, size=5)
compute_reciprocals(values)
```

```
big_array = np.random.randint(1, 100, size=1000000)
%timeit compute_reciprocals(big_array)
```

```
1 loop, best of 3: 2.91 s per loop
```

# Vectorized Operations

- ▶ Repeated operations in loops are slow
- ▶ Package supports C extensions that allow for execution across arrays element by element
- ▶ Processing concurrently instead of sequentially saves time

```
print(compute_reciprocals(values))  
print(1.0 / values)
```

[ 0.16666667	1.	0.25	0.25	0.125	]
[ 0.16666667	1.	0.25	0.25	0.125	]

# Vectorized Operations

- ▶ Repeated operations in loops are slow
- ▶ Package supports C extensions that allow for execution across arrays element by element
- ▶ Processing concurrently instead of sequentially saves time

```
print(compute_reciprocals(values))  
print(1.0 / values)
```

```
[ 0.16666667  1.          0.25          0.25          0.125        ]  
[ 0.16666667  1.          0.25          0.25          0.125        ]
```

```
%timeit (1.0 / big_array)
```

```
100 loops, best of 3: 4.6 ms per loop
```



# Vectorized Operations

- ▶ Possible to create your own operations through the *numba* package
  - ▶ Write a function meant for numbers
  - ▶ Decorate the function indicating data types
  - ▶ Call the function on arrays

```
from numba import vectorize, float64

@vectorize([float64(float64, float64)])
def f(x, y):
    return x + y
```

# Vectorized Operations

- ▶ Possible to create your own operations through the *numba* package
  - ▶ Write a function meant for numbers
  - ▶ Decorate the function indicating data types
  - ▶ Call the function on arrays

```
from numba import vectorize, float64

@vectorize([float64(float64, float64)])
def f(x, y):
    return x + y
```

```
>>> a = np.arange(6)
>>> f(a, a)
array([ 0,  2,  4,  6,  8, 10])
```

# Vectorized Operations

- ▶ Possible to create your own operations through the *numba* package
  - ▶ Write a function meant for numbers
  - ▶ Decorate the function indicating data types
  - ▶ Call the function on arrays

```
from numba import vectorize, float64

@vectorize([float64(float64, float64)])
def f(x, y):
    return x + y
```

```
>>> a = np.arange(6)
>>> f(a, a)
array([ 0,  2,  4,  6,  8, 10])
```

# Agenda

- ▶ Review

- ▶ numpy



- ▶ matplotlib

- ▶ Lesson

- ▶ notebooks and scripts

- ▶ pandas

- ▶ Readings

- ▶ Python for Data Analysis by Wes McKinney

- ▶ <http://pandas.pydata.org/pandas-docs/stable/index.html>

## Objectives

- ▶ What are different data types? How can we choose appropriate charts for data types?
- ▶ What are good practices for visualization
  - ▶ Conditioning
  - ▶ Scale
  - ▶ Jiggling Baseline
  - ▶ Transformation

# Agenda

- ▶ Review

- ▶ numpy

 matplotlib

- ▶ Lesson

- ▶ notebooks and scripts

- ▶ pandas

- ▶ Readings

- ▶ Python for Data Analysis by Wes McKinney

- ▶ <http://pandas.pydata.org/pandas-docs/stable/index.html>

## Objectives





- ▶ What are different data types? How can we choose appropriate charts for data types?
- ▶ What are good practices for visualization
  - ▶ Conditioning
  - ▶ Scale
  - ▶ Jiggling Baseline
  - ▶ Transformation

```
import matplotlib as mpl  
import matplotlib.pyplot as plt
```

# matplotlib

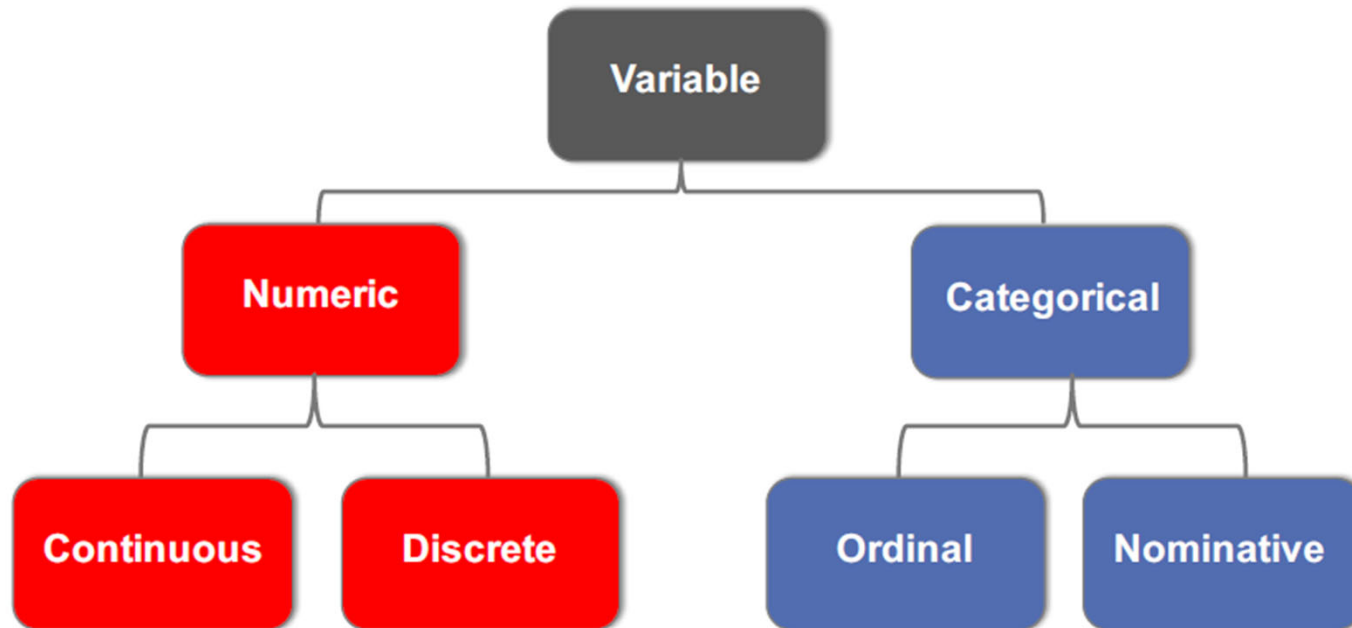
## ► Why visualization?

Property	Value	Accuracy
Mean of $x$	9	exact
Sample variance of $x$	11	exact
Mean of $y$	7.50	to 2 decimal places
Sample variance of $y$	4.125	$\pm 0.003$




# matplotlib

## ► Data Types



# matplotlib


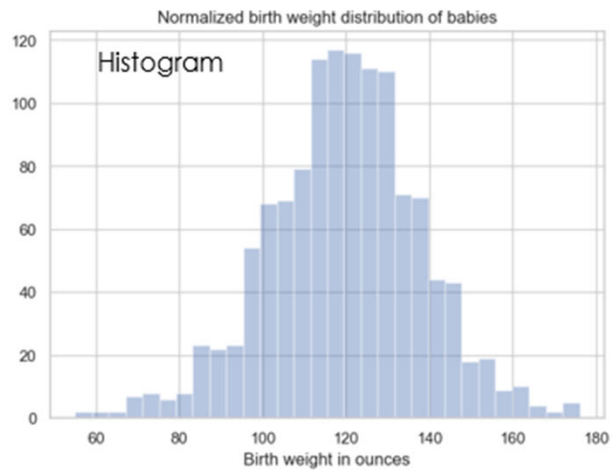
- One Quantitative Variable






# matplotlib

## ► One Quantitative Variable




# matplotlib

## ► One Quantitative Variable




# matplotlib

## ► One Qualitative Variable




# matplotlib

- One Qualitative Variable and One Quantitative Variable



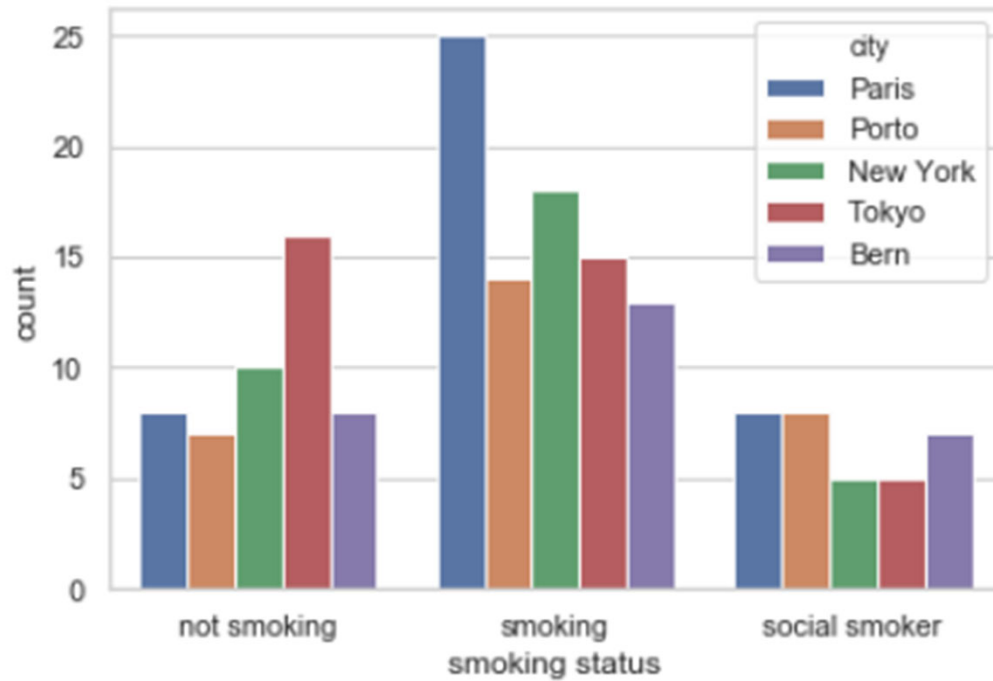
# matplotlib

- One Qualitative Variable and One Quantitative Variable




# matplotlib

## ► Multiple Qualitative Variables



# matplotlib

- ▶ Multiple Quantitative Variables
- ▶ Check out `sns.pairplot` !




# matplotlib

## ► Multiple Quantitative Variables







# matplotlib




matplotlib




# matplotlib




# matplotlib



# matplotlib



matplotlib



# Agenda

- ▶ Review

  - ▶ numpy

  - ▶ matplotlib

- ▶ Lesson

  -  notebooks and scripts

  - ▶ pandas

- ▶ Readings

  - ▶ Python for Data Analysis by Wes McKinney

  - ▶ <http://pandas.pydata.org/pandas-docs/stable/index.html>

## Objectives

- ▶ What are the advantages and disadvantages of scripts and notebooks?
- ▶ Is it possible to link notebook and scripts?
- ▶ What are the command line tools to convert between scripts and notebooks?

# Jupyter

- ▶ Notebooks

- ▶ Advantages

- ▶ Combine text, code and charts

- ▶ Support multiple languages including Python and R



# Jupyter

## ▶ Notebooks

### ▶ Advantages

- ▶ Combine text, code and charts
- ▶ Support multiple languages including Python and R

### ▶ Disadvantages

- ▶ JSON format makes version control difficult
- ▶ Lacking debugger to determine errors
- ▶ Refactoring difficult across multiple cells which may not reflect logical order of program

# Jupyter

## ► Notebooks

### ► Advantages


- Combine text, code and charts
- Support multiple languages including Python and R

### ► Disadvantages

- JSON format makes version control difficult
- Lacking debugger to determine errors
- Refactoring difficult across multiple cells which may not reflect logical order of program

*Link a notebook and a script so changes in one become changes in the other!*

# Agenda

- ▶ Review
  - ▶ numpy
  - ▶ matplotlib
- ▶ Lesson
  - ▶ notebooks and scripts
-  pandas
- ▶ Readings
  - ▶ Python for Data Analysis by Wes McKinney
  - ▶ <http://pandas.pydata.org/pandas-docs/stable/index.html>

## Objectives

- ▶ What are the components of pandas for storing data?
  - ▶ Series
  - ▶ DataFrame
- ▶ How can we access data in a table?

`import pandas as pd`

# pandas

- ▶ Tabular data consisting of rows and columns common in data analysis
  - ▶ Rows are observations in sample
  - ▶ Columns are features of the data
- ▶ Often *panel data* with rows consisting of timestamps

# pandas

- ▶ Tabular data consisting of rows and columns common in data analysis
  - ▶ Rows are observations in sample
  - ▶ Columns are features of the data
- ▶ Often *panel data* with rows consisting of timestamps
- ▶ pandas packages takes data structures and operations from languages like R and SQL
- ▶ Built on top of
  - ▶ numpy
  - ▶ scipy
  - ▶ some components of matplotlib

# pandas

- ▶ Tabular data consisting of rows and columns common in data analysis
  - ▶ Rows are observations in sample
  - ▶ Columns are features of the data
- ▶ Often *panel data* with rows consisting of timestamps

```
SELECT e.emp_id,  
       e.emp_name,  
       d.dept_name  
FROM Employee e  
INNER JOIN Department d ON e.dept_id = d.dept_id  
WHERE d.dept_name = 'finance'  
      AND e.emp_name LIKE '%A%'  
      AND e.salary > 500;
```

- ▶ pandas packages takes data structures and operations from languages like R and SQL
- ▶ Built on top of
  - ▶ numpy
  - ▶ scipy
  - ▶ some components of matplotlib

# Series

- ▶ One-dimensional object containing
  - ▶ Data
  - ▶ Labels (called *index*)
- ▶ Like array in numpy

```
# Creating a series  
index = ['a', 'b', 'c', 'd', 'e']  
series = pd.Series(np.arange(5), index=index)  
print(series)
```

```
a    0  
b    1  
c    2  
d    3  
e    4  
dtype: int64
```

# Series

- ▶ One-dimensional object containing
  - ▶ Data
  - ▶ Labels (called *index*)
  - ▶ Default index is range(N) where N is the length
- ▶ Index used for
  - ▶ lookups
  - ▶ aligning tables
  - ▶ supports hierarchical indexes, where each label is a tuple

```
# Creating a series  
index = ['a', 'b', 'c', 'd', 'e']  
series = pd.Series(np.arange(5), index=index)  
print(series)
```

```
a    0  
b    1  
c    2  
d    3  
e    4  
dtype: int64
```



# DataFrame

- ▶ Two-dimensional object containing
  - ▶ Data
  - ▶ Labels (called *index*)
  - ▶ Columns (ordered)
- ▶ Like a dictionary of Series where all Series have the same index

```
# Creating a dataframe with a dictionary
d = {'state' : ['FL', 'FL', 'GA', 'GA', 'GA'],
     'year'  : [2010, 2011, 2008, 2010, 2011],
     'pop'   : [18.8, 19.1, 9.7, 9.7, 9.8]}

df_d = pd.DataFrame(d)
print(df_d)
```

	pop	state	year
0	18.8	FL	2010
1	19.1	FL	2011
2	9.7	GA	2008
3	9.7	GA	2010
4	9.8	GA	2011

# Input and Output

- ▶ We can store tabular data in many formats
  - ▶ Comma Separated Values (CSV)
  - ▶ Tab Separated Values (TSV)
- ▶ Note that these file formats are not nested
  - ▶ Each row and column contains one entry

```
# the first row becomes the column indices  
df = pd.read_csv('simple.csv')  
print(df)
```

```
print(df.columns.values)
```

	a	b	c	d	message
0	1	2	3	4	hello
1	5	6	7	8	world
2	9	10	11	12	foo

```
['a' 'b' 'c' 'd' 'message']
```

# Demo

## Take-Aways

- ▶ What is Series
- ▶ What is DataFrame
- ▶ How does pandas support operations like numpy?
- ▶ How does pandas support operations like query languages?

# Summary

- ▶ Review

- ▶ numpy
- ▶ matplotlib

- ▶ Demos

- ▶ notebooks and scripts
- ▶ pandas

- ▶ Readings

- ▶ [http://wiki.scipy.org/Tentative\\_NumPy\\_Tutorial](http://wiki.scipy.org/Tentative_NumPy_Tutorial)
- ▶ <http://matplotlib.org/Matplotlib.pdf>